

```
In [1]: import pandas as pd
import matplotlib as plt
import numpy as np
```

```
In [2]: filename = "auto.csv"
headers = ["symboling", "normalized-losses", "make", "fuel-type",
           "aspiration", "num-of-doors", "body-style",
           "drive-wheels", "engine-location", "wheel-base",
           "length", "width", "height", "curb-weight", "engine-type",
           "num-of-cylinders", "engine-size", "fuel-system",
           "bore", "stroke", "compression-ratio", "horsepower",
           "peak-rpm", "city-mpg", "highway-mpg", "price"]
df=pd.read_csv(filename, names=headers)
```

```
In [3]: df.head(10)
```

```
Out[3]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...
5	2	?	audi	gas	std	two	sedan	fwd	front	99.8	...
6	1	158	audi	gas	std	four	sedan	fwd	front	105.8	...
7	1	?	audi	gas	std	four	wagon	fwd	front	105.8	...
8	1	158	audi	gas	turbo	four	sedan	fwd	front	105.8	...
9	0	?	audi	gas	turbo	two	hatchback	4wd	front	99.5	...

10 rows × 26 columns



Clean missing Data

```
In [4]: df.replace("?", np.nan, inplace = True)
df.head(10)
```

```
Out[4]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...
0	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...
1	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...
2	1	NaN	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...
5	2	NaN	audi	gas	std	two	sedan	fwd	front	99.8	...
6	1	158	audi	gas	std	four	sedan	fwd	front	105.8	...
7	1	NaN	audi	gas	std	four	wagon	fwd	front	105.8	...
8	1	158	audi	gas	turbo	four	sedan	fwd	front	105.8	...
9	0	NaN	audi	gas	turbo	two	hatchback	4wd	front	99.5	...

10 rows × 26 columns



```
In [5]: df.isnull().sum().sort_values(ascending=False)
```

```
Out[5]: normalized-losses    41
price                      4
stroke                     4
bore                       4
peak-rpm                   2
horsepower                 2
num-of-doors               2
length                     0
make                       0
fuel-type                  0
aspiration                 0
body-style                 0
drive-wheels               0
engine-location            0
wheel-base                0
height                    0
width                     0
highway-mpg               0
curb-weight               0
engine-type               0
num-of-cylinders          0
engine-size               0
fuel-system               0
compression-ratio         0
city-mpg                  0
symboling                 0
dtype: int64
```

```
In [6]: avg_n1 = df['normalized-losses'].astype(np.float64).mean(axis=0)

print('normalized-losses:', avg_n1)

df['normalized-losses'].replace(np.nan, avg_n1, inplace=True)
df['normalized-losses'] = df['normalized-losses'].astype(np.int64)
```

normalized-losses: 122.0

```
In [7]: avg_bore = df['bore'].astype(np.float64).mean(axis=0)

print('bore:', avg_bore)

df['bore'].replace(np.nan, avg_bore, inplace=True)
df['bore'] = df['bore'].astype(np.float64)
```

bore: 3.3297512437810957

```
In [8]: avg_stroke = df['stroke'].astype(np.float64).mean(axis=0)

print('stroke:', avg_stroke)

df['stroke'].replace(np.nan, avg_stroke, inplace=True)
df['stroke'] = df['stroke'].astype(np.float64)
```

stroke: 3.2554228855721337

```
In [9]: avg_horsepower = df['horsepower'].astype(np.float64).mean(axis=0)

print('horsepower:', avg_horsepower)

df['horsepower'].replace(np.nan, avg_horsepower, inplace=True)
df['horsepower'] = df['horsepower'].astype(np.float64)
```

horsepower: 104.25615763546799

```
In [10]: avg_peak_rpm = df['peak-rpm'].astype(np.float64).mean(axis=0)

print('peak-rpm:', avg_peak_rpm)

df['peak-rpm'].replace(np.nan, avg_peak_rpm, inplace=True)
df['peak-rpm'] = df['peak-rpm'].astype(np.float64)
```

peak-rpm: 5125.369458128079

```
In [11]: df.isnull().sum().sort_values(ascending=False)
```

```
Out[11]: price                4
num-of-doors                2
highway-mpg                 0
normalized-losses           0
make                        0
fuel-type                   0
aspiration                   0
body-style                   0
drive-wheels                 0
engine-location              0
wheel-base                  0
length                      0
width                       0
height                      0
curb-weight                  0
engine-type                  0
num-of-cylinders             0
engine-size                  0
fuel-system                  0
bore                         0
stroke                      0
compression-ratio            0
horsepower                   0
```

```
peak-rpm      0
city-mpg      0
symboling     0
dtype: int64
```

```
In [12]: print(df["num-of-doors"].value_counts())
```

```
four    114
two      89
Name: num-of-doors, dtype: int64
```

```
In [13]: df["num-of-doors"].replace(np.nan, "four", inplace=True)
```

```
In [14]: df.isnull().sum().sort_values(ascending=False)
```

```
Out[14]: price      4
highway-mpg      0
normalized-losses  0
make             0
fuel-type        0
aspiration       0
num-of-doors     0
body-style       0
drive-wheels     0
engine-location  0
wheel-base      0
length          0
width           0
height          0
curb-weight     0
engine-type     0
num-of-cylinders 0
engine-size     0
fuel-system     0
bore            0
stroke         0
compression-ratio 0
horsepower     0
peak-rpm       0
city-mpg       0
symboling      0
dtype: int64
```

```
In [15]: df.duplicated().sum()
```

```
Out[15]: 0
```

```
In [16]: df.dropna(subset=["price"], axis=0, inplace=True)
df.reset_index(drop=True, inplace=True)
```

```
In [17]: df.isnull().sum().sort_values(ascending=False)
```

```
Out[17]: price      0
highway-mpg      0
normalized-losses  0
make             0
fuel-type        0
aspiration       0
num-of-doors     0
body-style       0
drive-wheels     0
engine-location  0
```

```

wheel-base      0
length          0
width           0
height          0
curb-weight     0
engine-type     0
num-of-cylinders 0
engine-size     0
fuel-system     0
bore            0
stroke         0
compression-ratio 0
horsepower     0
peak-rpm       0
city-mpg       0
symboling      0
dtype: int64

```

```
In [18]: df.shape
```

```
Out[18]: (201, 26)
```

```
In [19]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 201 entries, 0 to 200
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              201 non-null   int64
1   normalized-losses     201 non-null   int64
2   make                  201 non-null   object
3   fuel-type              201 non-null   object
4   aspiration             201 non-null   object
5   num-of-doors           201 non-null   object
6   body-style             201 non-null   object
7   drive-wheels           201 non-null   object
8   engine-location        201 non-null   object
9   wheel-base             201 non-null   float64
10  length                 201 non-null   float64
11  width                  201 non-null   float64
12  height                 201 non-null   float64
13  curb-weight            201 non-null   int64
14  engine-type            201 non-null   object
15  num-of-cylinders       201 non-null   object
16  engine-size            201 non-null   int64
17  fuel-system            201 non-null   object
18  bore                   201 non-null   float64
19  stroke                 201 non-null   float64
20  compression-ratio      201 non-null   float64
21  horsepower              201 non-null   float64
22  peak-rpm               201 non-null   float64
23  city-mpg               201 non-null   int64
24  highway-mpg            201 non-null   int64
25  price                  201 non-null   object
dtypes: float64(9), int64(6), object(11)
memory usage: 41.0+ KB

```

```
In [20]: df['diesel'] = np.where(df['fuel-type']=='diesel',1,0)
df['diesel'] = df['diesel'].astype(np.int64)
```

```
In [21]: df['gas'] = np.where(df['fuel-type']=='gas',1,0)
df['gas'] = df['gas'].astype(np.int64)
```

```
In [22]: df['price'] = df['price'].astype(np.int64)
```

```
In [23]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 201 entries, 0 to 200
Data columns (total 28 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              201 non-null    int64
1   normalized-losses      201 non-null    int64
2   make                   201 non-null    object
3   fuel-type              201 non-null    object
4   aspiration              201 non-null    object
5   num-of-doors           201 non-null    object
6   body-style             201 non-null    object
7   drive-wheels           201 non-null    object
8   engine-location        201 non-null    object
9   wheel-base             201 non-null    float64
10  length                 201 non-null    float64
11  width                  201 non-null    float64
12  height                 201 non-null    float64
13  curb-weight            201 non-null    int64
14  engine-type            201 non-null    object
15  num-of-cylinders       201 non-null    object
16  engine-size            201 non-null    int64
17  fuel-system            201 non-null    object
18  bore                   201 non-null    float64
19  stroke                 201 non-null    float64
20  compression-ratio      201 non-null    float64
21  horsepower             201 non-null    float64
22  peak-rpm               201 non-null    float64
23  city-mpg               201 non-null    int64
24  highway-mpg            201 non-null    int64
25  price                  201 non-null    int64
26  diesel                 201 non-null    int64
27  gas                    201 non-null    int64
dtypes: float64(9), int64(9), object(10)
memory usage: 44.1+ KB
```

Data Normalization

```
In [24]: lwh = df[['length', 'width', 'height']]
lwh_min=lwh.min()
lwh_max=lwh.max()
lwh_norm = (lwh-lwh_min)/(lwh_max-lwh_min)
lwh = lwh_norm
lwh.head(201)
```

```
Out[24]:
```

	length	width	height
0	0.413433	0.324786	0.083333
1	0.413433	0.324786	0.083333
2	0.449254	0.444444	0.383333
3	0.529851	0.504274	0.541667
4	0.529851	0.521368	0.541667
...

	length	width	height
196	0.711940	0.735043	0.641667
197	0.711940	0.726496	0.641667
198	0.711940	0.735043	0.641667
199	0.711940	0.735043	0.641667
200	0.711940	0.735043	0.641667

201 rows × 3 columns

```
In [25]: from scipy import stats
```

```
In [26]: coef_wheel_base, p_value_wheel_base = stats.pearsonr(df['wheel-base'],
                                                         df['price'])
coef_hp, p_value_hp = stats.pearsonr(df['horsepower'], df['price'])
coef_length, p_value_length = stats.pearsonr(df['length'], df['price'])
coef_width, p_value_width = stats.pearsonr(df['width'], df['price'])
coef_height, p_value_height = stats.pearsonr(df['height'], df['price'])
coef_curb_weight, p_value_curb_weight = stats.pearsonr(df['curb-weight'],
                                                         df['price'])
coef_engine_size, p_value_engine_size = stats.pearsonr(df['engine-size'],
                                                         df['price'])
coef_bore, p_value_bore = stats.pearsonr(df['bore'], df['price'])
coef_city_mpg, p_value_city_mpg = stats.pearsonr(df['city-mpg'],
                                                  df['price'])
coef_highway_mpg, p_value_highway_mpg = stats.pearsonr(df['highway-mpg'],
                                                         df['price'])
coef_symboling, p_value_symboling = stats.pearsonr(df['symboling'],
                                                     df['price'])
coef_normalized_losses, p_value_normalized_losses = stats.pearsonr(df['normalized-losse
                                                         df['price'])
coef_diesel, p_value_diesel = stats.pearsonr(df['diesel'], df['price'])
coef_gas, p_value_gas = stats.pearsonr(df['gas'], df['price'])
coef_stroke, p_value_stroke = stats.pearsonr(df['stroke'], df['price'])
coef_compression_ratio, p_value_compression_ratio = stats.pearsonr(df['compression-rati
                                                         df['price'])
coef_peak_rpm, p_value_peak_rpm = stats.pearsonr(df['peak-rpm'],
                                                  df['price'])

pd.DataFrame({'': ['wheel-base vs price', 'horsepower vs price',
                  'length vs price', 'width vs price', 'height vs price',
                  'curb-weight vs price', 'engine-size vs price',
                  'bore vs price', 'city-mpg vs price',
                  'highway-mpg vs price', 'symboling vs price',
                  'normalized-losses vs price',
                  'diesel vs price',
                  'gas vs price', 'stroke vs price',
                  'compression-ratio vs price', 'peak-rpm vs price'],
             'Pearson Correlation Coefficient': [coef_wheel_base, coef_hp,
                                                  coef_length, coef_width,
                                                  coef_height,
                                                  coef_curb_weight,
                                                  coef_engine_size, coef_bore,
                                                  coef_city_mpg,
                                                  coef_highway_mpg,
                                                  coef_symboling,
```

```

coef_normalized_losses,
coef_diesel,coef_gas,
coef_stroke,
coef_compression_ratio,
coef_peak_rpm],
'P-value':[p_value_wheel_base,p_value_hp,p_value_length,
p_value_width,p_value_height,p_value_curb_weight,
p_value_engine_size,p_value_bore,
p_value_city_mpg,p_value_highway_mpg,
p_value_symboling,p_value_normalized_losses,
p_value_diesel,p_value_gas,
p_value_stroke,p_value_compression_ratio,
p_value_peak_rpm],
'Significant of correlation':['Yes','Yes','Yes','Yes','Yes',
'Yes','Yes','Yes','Yes','Yes',
'No','Yes','No',
'No','No','No','No']

```

```

}))

```

Out[26]:

		Pearson Correlation Coefficient	P-value	Significant of correlation
0	wheel-base vs price	0.584642	8.076488e-20	Yes
1	horsepower vs price	0.809575	6.369057e-48	Yes
2	length vs price	0.690628	8.016477e-30	Yes
3	width vs price	0.751265	9.200336e-38	Yes
4	height vs price	0.135486	5.514627e-02	Yes
5	curb-weight vs price	0.834415	2.189577e-53	Yes
6	engine-size vs price	0.872335	9.265492e-64	Yes
7	bore vs price	0.543155	8.049189e-17	Yes
8	city-mpg vs price	-0.686571	2.321132e-29	Yes
9	highway-mpg vs price	-0.704692	1.749547e-31	Yes
10	symboling vs price	-0.082391	2.449149e-01	No
11	normalized-losses vs price	0.133999	5.789749e-02	Yes
12	diesel vs price	0.110326	1.189625e-01	No
13	gas vs price	-0.110326	1.189625e-01	No
14	stroke vs price	0.082269	2.456161e-01	No
15	compression-ratio vs price	0.071107	3.158110e-01	No
16	peak-rpm vs price	-0.101616	1.511769e-01	No

In [27]:

```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.model_selection import train_test_split

regr = AdaBoostRegressor(random_state=0, n_estimators=100)

```



```

X=df[['highway-mpg','horsepower','curb-weight',
      'engine-size','wheel-base',
      'normalized-losses',
      'length','width','bore',
      'city-mpg','symboling','height','bore',
      'stroke','diesel']]
y=df['price']

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=1,
                                                    test_size=0.33)

clf = regr.fit(X_train, y_train)
clf.predict(X_test)
clf.score(X_test, y_test)

```

Out[27]: 0.8923459285852271

```

In [28]: from sklearn.ensemble import GradientBoostingRegressor

X=df[['highway-mpg','horsepower','curb-weight',
      'engine-size','wheel-base',
      'normalized-losses',
      'length','width','bore',
      'city-mpg','symboling','height','bore',
      'stroke','diesel']]
y=df['price']

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=1,
                                                    test_size=0.33)

clf = GradientBoostingRegressor().fit(X_train, y_train)
clf.predict(X_test)
clf.score(X_test, y_test)

```

Out[28]: 0.9012976558224672

```

In [37]: from sklearn.neural_network import MLPRegressor
from sklearn.model_selection import train_test_split

X=df[['highway-mpg','horsepower','curb-weight',
      'engine-size','wheel-base',
      'normalized-losses',
      'length','width','bore',
      'city-mpg']]
y=df['price']

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=1,
                                                    test_size=0.33)

regr = MLPRegressor(hidden_layer_sizes=(1000,1000), activation='relu',
                    solver='adam', alpha=0.001, batch_size='auto',
                    learning_rate='constant', learning_rate_init=0.01, power_t=0.5,
                    max_iter=1000, shuffle=True,
                    random_state=0, tol=0.0001, verbose=False, warm_start=False,
                    momentum=0.9, nesterovs_momentum=True,
                    early_stopping=False, validation_fraction=0.1, beta_1=0.9,
                    beta_2=0.999, epsilon=1e-08).fit(X_train, y_train)

regr.predict(X_test)
regr.score(X_test, y_test)

```

Out[37]: 0.8706127983267664

In []: